

# TrachtPS4

February 6, 2019

## 1 Problem Set 4

1.0.1 by Daniel Tracht, February 2019

### 1.1 Problem 1

For this problem, we use the 200 data points provided in the incomes.txt file, containing incomes reported in U.S. dollars. We will be using the log normal distribution

#### 1.1.1 Part a

We wish to plot a histogram of the percentages from our data with 30 bins:

```
In [1]: # Import the necessary libraries
import numpy as np
import scipy.stats as sts
import requests

# Download and save the data file incomes.txt
url = ('https://raw.githubusercontent.com/dtracht/persp-model-econ_W19/' +
       'master/ProblemSets/PS4/data/incomes.txt')
data_file = requests.get(url, allow_redirects=True)
open('data/incomes.txt', 'wb').write(data_file.content)

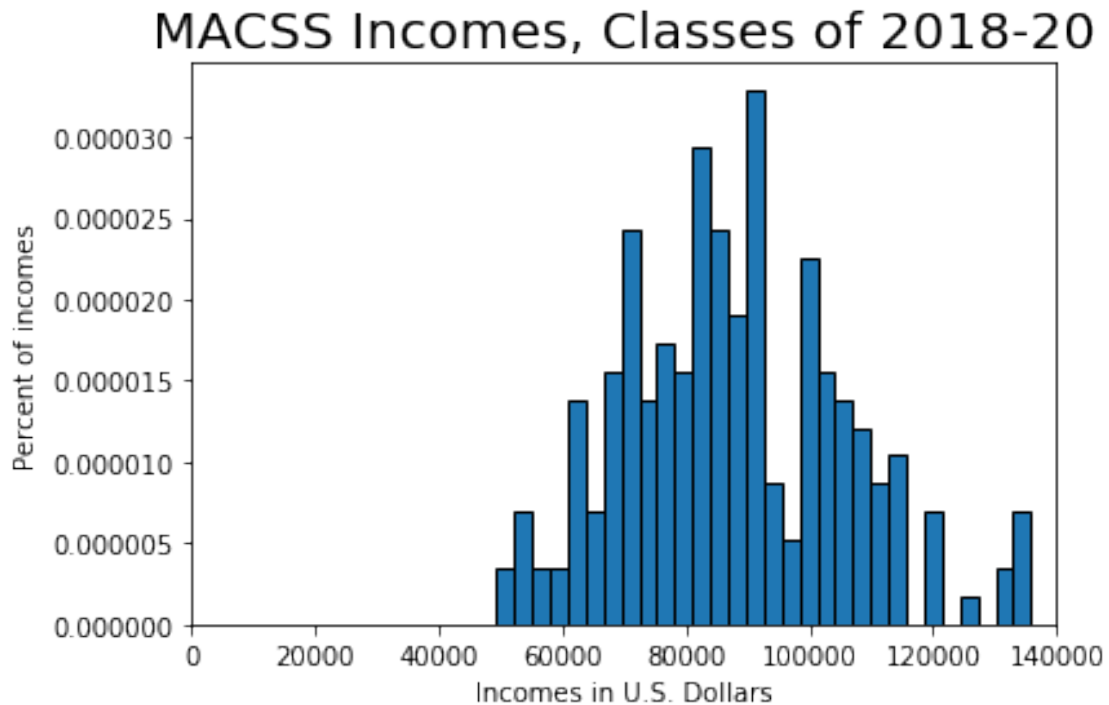
# Load the data as a NumPy array
pts = np.loadtxt('data/incomes.txt')

In [3]: import matplotlib.pyplot as plt

num_bins = 30
# normed option has been deprecated
# using density=True instead
count, bins, ignored = plt.hist(pts, num_bins, density=True,
                                edgecolor='k')

plt.title("MACSS Incomes, Classes of 2018-20", fontsize=20)
plt.xlabel("Incomes in U.S. Dollars")
plt.ylabel("Percent of incomes")
plt.xlim([0,140000])
```

Out [3]: (0, 140000)



### 1.1.2 Part b

We wish to plot the probability density function of the lognormal distribution,

$$f(x|\mu = 11.0, \sigma = 0.5)$$

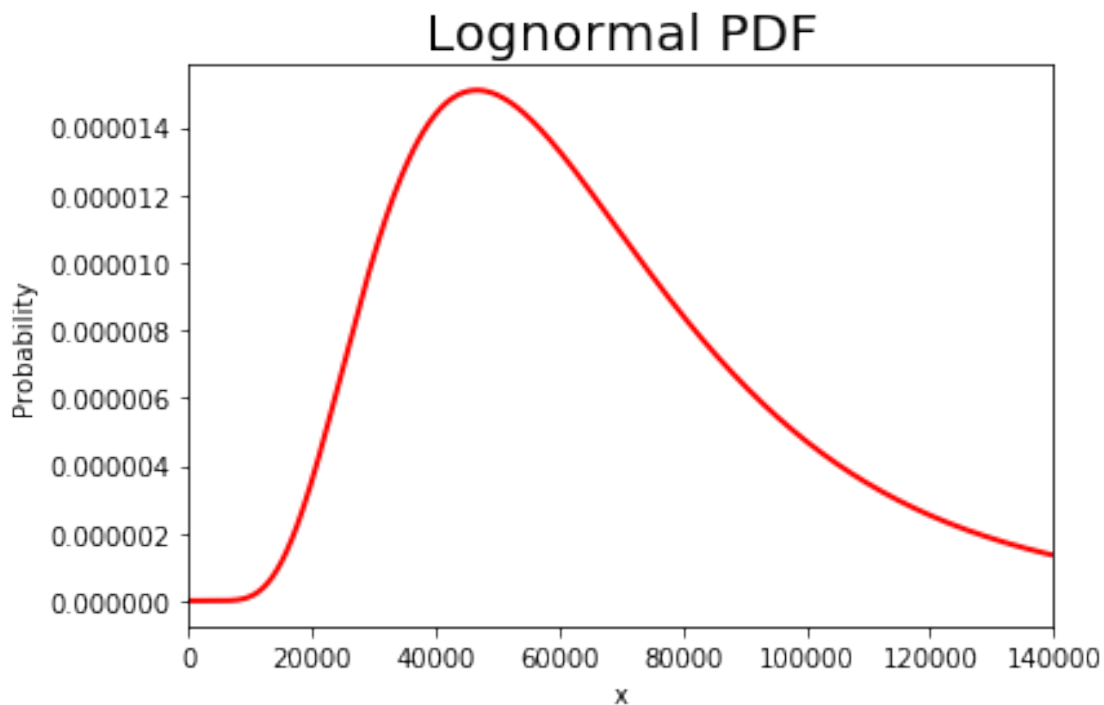
for

$$0 \leq x \leq 150000$$

:

```
In [4]: # takes vector of x values, parameters for mu and sigma
# returns vector of pdf values for the lognormal
from scipy.stats import lognorm
def lognormal_pdf (xvals, mu, sigma):
    pdf_vals = (1 / (xvals * sigma * np.sqrt(2 * np.pi))) * np.exp (-(np.log(xvals)-mu)**2 / (2 * sigma**2))
    #pdf_vals = lognorm.pdf(xvals, sigma, loc=mu, scale=np.exp(mu))
    #take pdf values that are rounded to zero and reset them to very low
    pdf_zero = pdf_vals==0
    pdf_vals[pdf_zero] = 1e-16
    return pdf_vals
```

```
In [5]: # cannot actually start at zero
dist_pts = np.linspace(1e-16, 150000, 1000)
mu_1 = 11
sig_1 = 0.5
plt.plot(dist_pts, lognormal_pdf(dist_pts, mu_1, sig_1),
         linewidth=2, color='r', label='1:  $\mu=11.0, \sigma=0.5$ ')
plt.title("Lognormal PDF", fontsize=20)
plt.xlabel("x")
plt.ylabel("Probability")
plt.xlim([0, 140000])
plt.show()
```




```
In [6]: # Plot histogram with the PDF overlaid
num_bins = 30
# normed option has been deprecated
# using density=True instead
count, bins, ignored = plt.hist(pts, num_bins, density=True,
                                edgecolor='k')
plt.title("MACSS Incomes, Classes of 2018-20", fontsize=20)
plt.xlabel("Incomes in U.S. Dollars")
plt.ylabel("Percent of incomes")
plt.xlim([0, 140000])

plt.plot(dist_pts, lognormal_pdf(dist_pts, mu_1, sig_1),
```

```

        linewidth=2, color='r', label='1:  $\mu=11.0, \sigma=0.5$ ')
plt.legend(loc='upper left')
plt.show()

```



We wish to find the log likelihood value for this particular parameterization of this function form and the given data:

```

In [7]: # function that takes a vector of x values, and parameters mu and sigma
        # returns the log likelihood for that vector for the lognormal distribution
def log_likelihood(xvals, mu, sigma):
    pdf_vals = lognormal_pdf(xvals, mu, sigma)
    log_pdf_vals = np.log(pdf_vals)
    log_likelihood = log_pdf_vals.sum()
    return log_likelihood

```

```

In [8]: print('Log-likelihood:', log_likelihood(pts, mu_1, sig_1))

```

Log-likelihood: -2385.856997808558

### 1.1.3 Part c

We wish to estimate the parameters for a log normal distribution to fit this data by the method of maximum likelihood estimation:

```

In [9]: # a critereon function to be passed to the minimizer
def crit(params, *args):
    mu, sigma = params
    xvals, junk = args
    log_likelihood_value = log_likelihood(xvals, mu, sigma)
    neg_log_likelihood_value = -log_likelihood_value

    return neg_log_likelihood_value

In [10]: import scipy.optimize as opt

    # starting from our values in part b
    mu_init = 11
    sig_init = 0.5
    params_init = np.array([mu_init, sig_init])
    mle_args = (pts, 0)
    # constraining sigma to be positive
    bnds = ((None, None), (1e-16, None))
    results = opt.minimize(crit, params_init, args=(mle_args), method='SLSQP', bounds=bnds)
    mu_MLE, sig_MLE = results.x

```

We wish to plot our new distribution as well as our original guess and the data:

```

In [11]: # Plot histogram with the old PDF and new PDF overlaid
num_bins = 30
# normed option has been deprecated
# using density=True instead
count, bins, ignored = plt.hist(pts, num_bins, density=True,
                                edgecolor='k')

plt.title("MACSS Incomes, Classes of 2018-20", fontsize=20)
plt.xlabel("Incomes in U.S. Dollars")
plt.ylabel("Percent of incomes")
plt.xlim([0,140000])


plt.plot(dist_pts, lognormal_pdf(dist_pts, mu_1, sig_1),
         linewidth=2, color='r', label='1:  $\mu=11.0, \sigma=0.5$ ')
plt.legend(loc='upper left')

plt.plot(dist_pts, lognormal_pdf(dist_pts, mu_MLE, sig_MLE),
         linewidth=2, color='g', label='MLE:  $\mu=11.4, \sigma=0.2$ ')
plt.legend(loc='upper left')

plt.show()

```

## MACSS Incomes, Classes of 2018-20



We wish to get the variance-covariance matrix for our estimates. However, using the constrained method above, the optimizer does not return an inverse Hessian object. To get it, we can simply use a method that returns one beginning at the solution to the constrained problem.

```
In [12]: # starting from our values in the constrained case
mu_init = mu_MLE
sig_init = sig_MLE
params_init = np.array([mu_init, sig_init])
mle_args = (pts, 0)

results = opt.minimize(crit, params_init, args=(mle_args))
mu_MLE_2, sig_MLE_2 = results.x

print('mu_MLE=', mu_MLE_2, ' sig_MLE=', sig_MLE_2)
print('MLE Log-likelihood:', log_likelihood(pts, mu_MLE_2, sig_MLE_2))

vcv_mle = results.hess_inv

stderr_mu_mle = np.sqrt(vcv_mle[0,0])
stderr_sig_mle = np.sqrt(vcv_mle[1,1])
print('VCV(MLE) = ', vcv_mle)
print('Standard error for mu estimate = ', stderr_mu_mle)
print('Standard error for sigma estimate = ', stderr_sig_mle)

mu_MLE= 11.359022994120751  sig_MLE= 0.20817731910735596
MLE Log-likelihood: -2241.7193013573587
```

```
VCV(MLE) = [[1.94308215e-04 1.27192004e-05]
 [1.27192004e-05 1.24567598e-04]]
Standard error for mu estimate = 0.013939448157069998
Standard error for sigma estimate = 0.01116098553849848
```

#### 1.1.4 Part d

We wish to perform a likelihood ratio test to determine the probability that the data came from the distribution using the parameters from part b:

```
In [13]: # null hypothesis coming from the parameters in part b
mu_new, sig_new = np.array([11, 0.5])
log_lik_h0 = log_likelihood(pts, mu_new, sig_new)
print('hypothesis value log likelihood', log_lik_h0)
log_lik_mle = log_likelihood(pts, mu_MLE, sig_MLE)
print('MLE log likelihood', log_lik_mle)
LR_val = 2 * (log_lik_mle - log_lik_h0)
print('likelihood ratio value', LR_val)
pval_h0 = 1.0 - sts.chi2.cdf(LR_val, 2)
print('chi squared of H0 with 2 degrees of freedom p-value = ', pval_h0)
```

```
hypothesis value log likelihood -2385.856997808558
MLE log likelihood -2241.7193013574033
likelihood ratio value 288.27539290230925
chi squared of H0 with 2 degrees of freedom p-value = 0.0
```

The probability of the data being observed under the assumption of the functional form and parameters in part b is so small it rounds to zero in floating point operation.

#### 1.1.5 Part e

We wish to estimate the probability that a graduate of MACSS will earn more than 100,000 dollars and the probability that a graduate of MACSS will earn less than 75,000 dollars using the model and parameters we have estimated. For this we can go to the CDF of the log normal distribution

```
In [14]: prob_under75 = lognorm.cdf(75000, sig_MLE, loc=mu_MLE, scale=np.exp(mu_MLE))
prob_over100 = 1-lognorm.cdf(100000, sig_MLE, loc=mu_MLE, scale=np.exp(mu_MLE))
print("Probability of earning less than $75,000:", prob_under75)
print("Probability of earning more than $100,000:", prob_over100)
```

```
Probability of earning less than $75,000: 0.2599983218978105
Probability of earning more than $100,000: 0.23003229599994834
```

## 1.2 Problem 2

In this problem we will be solving a linear regression using maximum likelihood estimation. Our model is

$$sick_i = \beta_0 + \beta_1 age_i + \beta_2 children_i + \beta_3 tempwinter_i + \varepsilon_i$$

where

$$\varepsilon_i \sim N(0, \sigma^2)$$

### 1.2.1 Part a

We begin by estimating the parameters through maximum likelihood estimation:

```
In [15]: # Download and save the data file incomes.txt
url = ('https://raw.githubusercontent.com/dtracht/persp-model-econ_W19/' +
       'master/ProblemSets/PS4/data/sick.txt')
data_file = requests.get(url, allow_redirects=True)
open('data/sick.txt', 'wb').write(data_file.content)

# Load the data as a NumPy array
pts = np.genfromtxt("data/sick.txt", delimiter=",", skip_header=1)

y_pts = pts[:,0]
x1_pts = pts[:,1]
x2_pts = pts[:,2]
x3_pts = pts[:,3]

In [16]: # takes vector of x values, sigma
# returns vector of pdf values for the normal with mean 0
from scipy.stats import norm
def normal_pdf (xvals, sigma):
    # reset sigma to absolute value
    # get around negative issue here as opposed to solving and then getting the VCV m
    # don't see a way around this now, but it feels dirty
    sigma = abs(sigma)
    pdf_vals = (1 / np.sqrt(2 * np.pi * sigma**2)) * np.exp(-(xvals) ** 2 / (2 * sigma
    #pdf_vals = norm.pdf(xvals, loc=0, scale=sigma)
    # take values rounded to zero and reset to very low
    pdf_zero = pdf_vals == 0
    pdf_vals[pdf_zero] = 1e-16
    return pdf_vals

In [17]: # function that takes an array of data, vector of betas, and sigma parameter
# returns the log likelihood for that vector for the normal distribution with mean 0
def log_likelihood_normal(beta_0, beta_1, beta_2, beta_3, sigma, y_vals, x1_vals, x2_
    # calculates epsilon values for beta vector and data
    epsilon_vals = y_vals - beta_0 - beta_1*x1_vals - beta_2*x2_vals - beta_3*x3_vals
    pdf_vals = normal_pdf(epsilon_vals, sigma)
    log_pdf_vals = np.log(pdf_vals)
```



```

log_likelihood = log_pdf_vals.sum()
return log_likelihood

```

In [18]: *# a critereon function to be passed to the minimizer*

```

def crit_normal(params, *args):
    beta_0, beta_1, beta_2, beta_3, sigma = params
    y_vals, x1_vals, x2_vals, x3_vals = args
    log_likelihood_value = log_likelihood_normal(beta_0, beta_1, beta_2, beta_3, sigma)
    neg_log_likelihood_value = -log_likelihood_value
    return neg_log_likelihood_value

```

In [19]: *import scipy.optimize as opt*

```

# starting from visual inspection
beta_0_init = 0
beta_1_init = 0
beta_2_init = 0
beta_3_init = 0
sigma_init = 1
params_init = np.array([beta_0_init, beta_1_init, beta_2_init, beta_3_init, sigma_init])
mle_args = (y_pts, x1_pts, x2_pts, x3_pts)

# constraining sigma to be positive
#bnds = ((None, None), (None, None), (None, None), (None, None), (1e-16, None))
#results = opt.minimize(crit_normal, params_init, args=(mle_args), method='L-BFGS-B',
results = opt.minimize(crit_normal, params_init, args=(mle_args))
beta_0_MLE, beta_1_MLE, beta_2_MLE, beta_3_MLE, sigma_MLE = results.x

print('beta_0_MLE=', beta_0_MLE)
print('beta_1_MLE=', beta_1_MLE)
print("beta_2_MLE=", beta_2_MLE)
print("beta_3_MLE=", beta_3_MLE)
print("sigma_MLE=", sigma_MLE)
print('MLE Log-likelihood:', log_likelihood_normal(beta_0_MLE, beta_1_MLE, beta_2_MLE, beta_3_MLE, sigma_MLE))

vcv_mle = results.hess_inv
print('VCV(MLE) = ', vcv_mle)

```

```

beta_0_MLE= 0.25164638362713215
beta_1_MLE= 0.012933350045131211
beta_2_MLE= 0.400502048301034
beta_3_MLE= -0.009991673036254978
sigma_MLE= 0.003017682175869847
MLE Log-likelihood: 876.8650462887678
VCV(MLE) = [[ 9.21562214e-08 -6.95529550e-10 -5.02282376e-09 -9.62673307e-10
 -2.62646033e-08]
 [-6.95529550e-10  6.61646149e-10 -2.29838661e-09 -4.87517749e-10
  1.64799761e-09]

```

```

[-5.02282376e-09 -2.29838661e-09  2.31900554e-08  1.25339295e-09
 -1.30889427e-08]
[-9.62673307e-10 -4.87517749e-10  1.25339295e-09  4.29373073e-10
 -4.75934351e-10]
[-2.62646033e-08  1.64799761e-09 -1.30889427e-08 -4.75934351e-10
 1.66995820e-08]]

```

```

In [ ]: '''
        # starting from our values in the constrained case
        beta_0_init_2 = beta_0_MLE
        beta_1_init_2 = beta_1_MLE
        beta_2_init_2 = beta_2_MLE
        beta_3_init_2 = beta_3_MLE
        sigma_init_2 = sigma_MLE

        params_init_2 = np.array([beta_0_init_2, beta_1_init_2, beta_2_init_2, beta_3_init_2,
        print(params_init_2)
        mle_args = (y_pts, x1_pts, x2_pts, x3_pts)

        results = opt.minimize(crit_normal, params_init_2, args=(mle_args))
        beta_0_MLE_2, beta_1_MLE_2, beta_2_MLE_2, beta_3_MLE_2, sigma_MLE_2 = results.x

        print('beta_0_MLE=', beta_0_MLE_2)
        print('beta_1_MLE=', beta_1_MLE_2)
        print("beta_2_MLE=", beta_2_MLE_2)
        print("beta_3_MLE=", beta_3_MLE_2)
        print("sigma_MLE=", sigma_MLE_2)
        print('MLE Log-likelihood:', log_likelihood_normal(beta_0_MLE_2, beta_1_MLE_2, beta_2_MLE_2,
        sigma_MLE_2))

        vcv_mle = results.hess_inv
        print('VCV(MLE) = ', vcv_mle)
        '''

```

## 1.2.2 Part b

We wish to use a likelihood ratio test to determine the probability that  $\beta_0 = 1.0$ ,  $\sigma^2 = 0.01$ , and  $\beta_1 = \beta_2 = \beta_3 = 0$ :

```

In [20]: beta_0_null = 1
        beta_1_null = 0
        beta_2_null = 0
        beta_3_null = 0
        sigma_null = 0.01
        log_lik_h0 = log_likelihood_normal(beta_0_null, beta_1_null, beta_2_null, beta_3_null,
        sigma_null)
        print('hypothesis value log likelihood', log_lik_h0)
        log_lik_mle = log_likelihood_normal(beta_0_MLE, beta_1_MLE, beta_2_MLE, beta_3_MLE,
        sigma_MLE)
        print('MLE log likelihood', log_lik_mle)

```

```
LR_val = 2 * (log_lik_mle - log_lik_h0)
print('likelihood ratio value', LR_val)
pval_h0 = 1.0 - sts.chi2.cdf(LR_val, 2)
print('chi squared of H0 with 2 degrees of freedom p-value = ', pval_h0)
```

```
hypothesis value log likelihood -30663.099086401762
MLE log likelihood 876.8650462887678
likelihood ratio value 63079.92826538106
chi squared of H0 with 2 degrees of freedom p-value = 0.0
```

We strongly reject the hypothesis that the data was generated with a true process of the given function and parameters.