

TrachtPS6

February 19, 2019

1 Problem Set 6

1.1 Daniel Tracht

1.2 Problem 1

1.2.1 Part a

We wish to import the data from the Auto.csv file, and replace the values that seem to be out of place.

```
In [1]: import pandas as pd
```

```
autos = pd.read_csv('data/Auto.csv')
```

```
autos.dtypes
```

```
Out[1]: mpg          float64
cylinders          int64
displacement       float64
horsepower         object
weight            int64
acceleration       float64
year              int64
origin            int64
name              object
dtype: object
```

We expect horsepower to be a numeric.

```
In [2]: autos.sort_values(by="horsepower").tail()
```

```
Out[2]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
336	23.6	4	140.0	?	2905	14.3	80	
126	21.0	6	200.0	?	2875	17.0	74	
354	34.5	4	100.0	?	2320	15.8	81	
32	25.0	4	98.0	?	2046	19.0	71	
330	40.9	4	85.0	?	1835	17.3	80	

	origin	name
336	1	ford mustang cobra
126	1	ford maverick
354	2	renault 18i
32	1	ford pinto
330	2	renault lecar deluxe

We have a question mark for our not applicable values. Let's reimport with the right option.

```
In [3]: autos = pd.read_csv('data/Auto.csv', na_values="?")
```

1.2.2 Part b

We wish to produce a scatterplot matrix which includes all of the quantatative variables:

```
In [5]: from pandas.plotting import scatter_matrix

# making a data frame of the quantative variables
# while the origin variable is stored as a numeric in our data, it is a categor
# autos is a pretty well known data set
# 1 is USA, 2 is Europe, 3 is Japan
#df_quant = autos[["mpg", "cylinders", "displacement", "horsepower", "weight",
#                  "acceleration", "year", "origin"]]
# ~ would be the line you asked for
df_quant = autos[["mpg", "cylinders", "displacement", "horsepower",
                  "weight", "acceleration", "year"]]

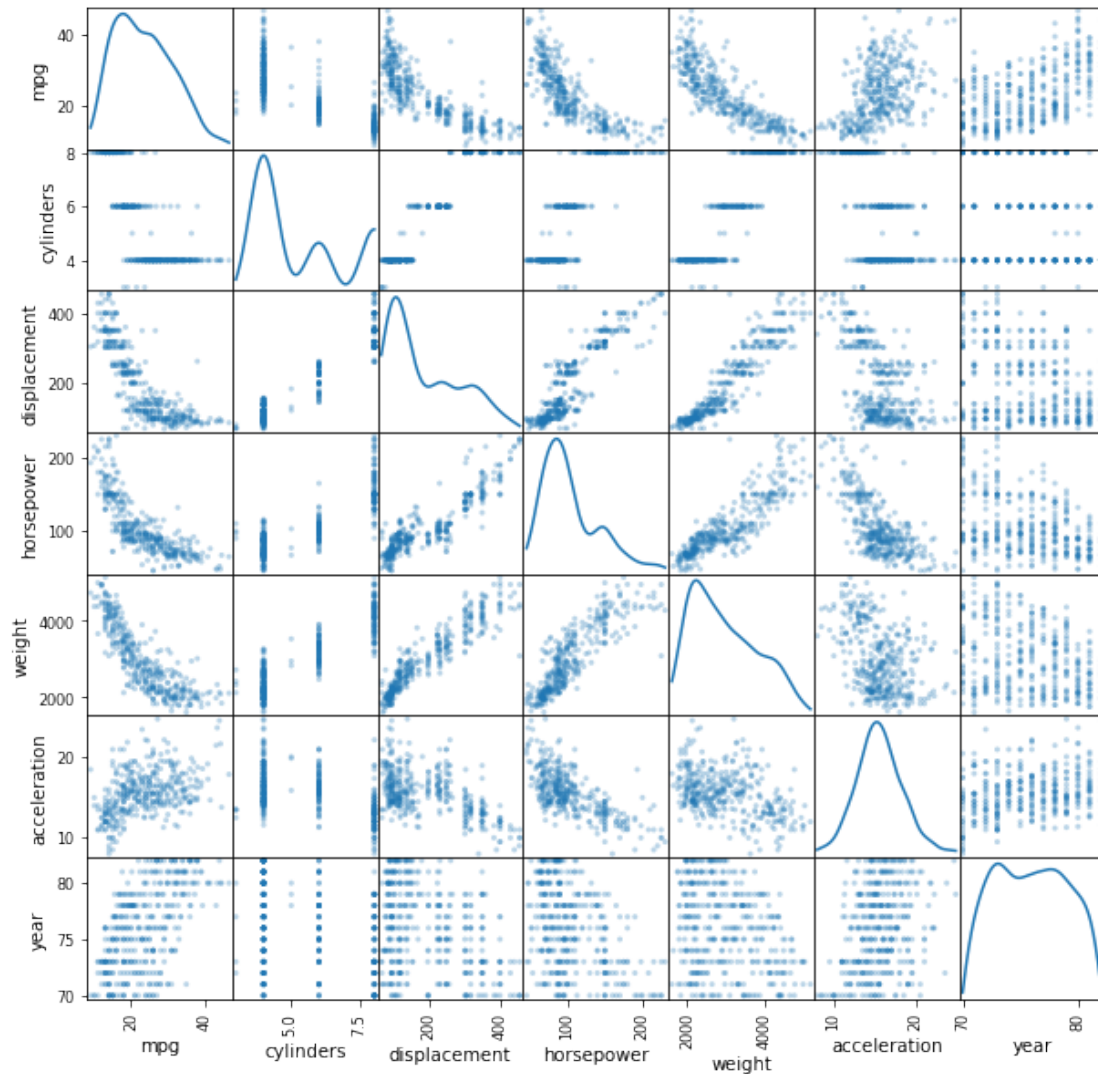
#scatter_matrix(df_quant, alpha=0.3, figsize=(6, 6), diagonal="kde")
# ~ would be the line you asked for, but 6, 6 was just too small to see anything
scatter_matrix(df_quant, alpha=0.3, figsize=(10, 10), diagonal="kde")

Out[5]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBAA86908>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFCDD710>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFD03C88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFD360B8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFD5D630>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFD84BA8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFDB6160>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFE1C710>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFE1C748>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFE78208>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFE9D780>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFEC6CF8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFF782B0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFF9D828>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFFC6DA0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BBFFF7358>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC00218D0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0048E48>],
```

```

<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC04D8400>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0502978>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0529EF0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC055B4A8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0581A20>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC05ABF98>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC05DB550>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0602AC8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0636080>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC065C5F8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0686B70>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC06B6128>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC06DF6A0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0706C18>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC07371D0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC075F748>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0787CC0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC07B7278>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC07DE7F0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0806D68>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0837320>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0860898>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC088AE10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC08B93C8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC08E1940>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC090BEB8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC0939470>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC09649E8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC098DF60>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC09BC518>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000020BC09E2A90>]],
dtype=object)

```



1.2.3 Part c

We wish to compute the correlation matrix for the quantitative variables:

In [6]: `df_quant.corr()`

```
Out [6]:
```

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.776260	-0.804443	-0.778427	-0.831739	
cylinders	-0.776260	1.000000	0.950920	0.842983	0.897017	
displacement	-0.804443	0.950920	1.000000	0.897257	0.933104	
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	
weight	-0.831739	0.897017	0.933104	0.864538	1.000000	
acceleration	0.422297	-0.504061	-0.544162	-0.689196	-0.419502	
year	0.581469	-0.346717	-0.369804	-0.416361	-0.307900	

	acceleration	year
mpg	0.422297	0.581469
cylinders	-0.504061	-0.346717
displacement	-0.544162	-0.369804
horsepower	-0.689196	-0.416361
weight	-0.419502	-0.307900
acceleration	1.000000	0.282901
year	0.282901	1.000000

1.2.4 Part d

We wish to estimate a multiple linear regression model:

```
In [7]: import statsmodels.api as sm
```

```
# defining a column of 1s as the constant
autos["const"] = 1

# making a dataframe of the exogenous variables

exog_origin = autos[["const", "cylinders", "displacement", "horsepower", "weight",
                    "acceleration", "year", "origin"]]
# ~ the line with origin as a quantative, not categorical variable

# We have to make some dummies first
origins = pd.get_dummies(autos["origin"], drop_first=True)
autos = pd.concat([autos, origins], axis=1)
autos.rename(columns={2: "Europe", 3: "Japan"}, inplace=True)

exog = autos[["const", "cylinders", "displacement", "horsepower", "weight",
              "acceleration", "year", "Europe", "Japan"]]

# running the regression
reg1_origin = sm.OLS(endog=autos['mpg'], exog=exog_origin, missing='drop')
results1_origin = reg1_origin.fit()
print(results1_origin.summary())

reg1 = sm.OLS(endog=autos['mpg'], exog=exog, missing='drop')
results1 = reg1.fit()
print(results1.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          mpg    R-squared:                0.821
Model:                  OLS    Adj. R-squared:           0.818
Method:                 Least Squares    F-statistic:         252.4
Date:                  Tue, 19 Feb 2019    Prob (F-statistic):    2.04e-139
```

Time: 15:26:27 Log-Likelihood: -1023.5
 No. Observations: 392 AIC: 2063.
 Df Residuals: 384 BIC: 2095.
 Df Model: 7
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-17.2184	4.644	-3.707	0.000	-26.350	-8.087
cylinders	-0.4934	0.323	-1.526	0.128	-1.129	0.142
displacement	0.0199	0.008	2.647	0.008	0.005	0.035
horsepower	-0.0170	0.014	-1.230	0.220	-0.044	0.010
weight	-0.0065	0.001	-9.929	0.000	-0.008	-0.005
acceleration	0.0806	0.099	0.815	0.415	-0.114	0.275
year	0.7508	0.051	14.729	0.000	0.651	0.851
origin	1.4261	0.278	5.127	0.000	0.879	1.973

Omnibus: 31.906 Durbin-Watson: 1.309
 Prob(Omnibus): 0.000 Jarque-Bera (JB): 53.100
 Skew: 0.529 Prob(JB): 2.95e-12
 Kurtosis: 4.460 Cond. No. 8.59e+04

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 8.59e+04. This might indicate that there are strong multicollinearity or other numerical problems.

OLS Regression Results

Dep. Variable: mpg R-squared: 0.824
 Model: OLS Adj. R-squared: 0.821
 Method: Least Squares F-statistic: 224.5
 Date: Tue, 19 Feb 2019 Prob (F-statistic): 1.79e-139
 Time: 15:26:27 Log-Likelihood: -1020.5
 No. Observations: 392 AIC: 2059.
 Df Residuals: 383 BIC: 2095.
 Df Model: 8
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-17.9546	4.677	-3.839	0.000	-27.150	-8.759
cylinders	-0.4897	0.321	-1.524	0.128	-1.121	0.142
displacement	0.0240	0.008	3.133	0.002	0.009	0.039
horsepower	-0.0182	0.014	-1.326	0.185	-0.045	0.009
weight	-0.0067	0.001	-10.243	0.000	-0.008	-0.005
acceleration	0.0791	0.098	0.805	0.421	-0.114	0.272
year	0.7770	0.052	15.005	0.000	0.675	0.879

Europe	2.6300	0.566	4.643	0.000	1.516	3.744
Japan	2.8532	0.553	5.162	0.000	1.766	3.940
=====						
Omnibus:		23.395	Durbin-Watson:			1.291
Prob(Omnibus):		0.000	Jarque-Bera (JB):			34.452
Skew:		0.444	Prob(JB):			3.30e-08
Kurtosis:		4.150	Cond. No.			8.70e+04
=====						

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.7e+04. This might indicate that there are strong multicollinearity or other numerical problems.

If origin is included in the regression as a categorical variable, we find that β_0 , β_2 , β_4 , β_6 , and β_7 are statistically significant at the 1% level. We find that β_1 , β_3 , and β_5 are not statistically significant at the 10% level. In words, an automobile model that is 1 year newer would have 0.7508 more miles per gallon, ceteris paribus.

If origin is included as a categorical variable, we find that β_0 , β_2 , β_4 , β_6 , as well as the two coefficients for the origin dummies (being from Europe or Japan relative to the United States) are statistically significant at the 1% level. We find that β_1 , β_3 , and β_5 are not statistically significant at the 10% level. In words, an automobile model that is 1 year newer would have 0.7770 more miles per gallon, ceteris paribus.

1.2.5 Part e

From the scatterplot, it seems that displacement, horsepower, and weight are most likely to have a non-linear relationship with mpg_i . We wish to estimate a linear regression with a squared term to these three as well as $acceleration_i$:

```
In [8]: # Generating the square terms
autos["disp_sq"] = autos["displacement"]**2
autos["horses_sq"] = autos["horsepower"]**2
autos["weight_sq"] = autos["weight"]**2
autos["accel_sq"] = autos["acceleration"]**2

# taking the square terms into a data frame and joining them with the others
exog_sq = autos[["disp_sq", "horses_sq", "weight_sq", "accel_sq"]]
exog2_origin = pd.concat([exog_origin, exog_sq], axis=1)
exog2 = pd.concat([exog, exog_sq], axis=1)

# running the regression with origin as quantatative
reg2_origin = sm.OLS(endog=autos['mpg'], exog=exog2_origin, missing='drop')
results2_origin = reg2_origin.fit()
print(results2_origin.summary())

# running the regression with origin as categorical
```

```
reg2 = sm.OLS(endog=autos['mpg'], exog=exog2, missing='drop')
results2 = reg2.fit()
print(results2.summary())
```

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.870			
Model:	OLS	Adj. R-squared:	0.866			
Method:	Least Squares	F-statistic:	230.2			
Date:	Tue, 19 Feb 2019	Prob (F-statistic):	1.75e-160			
Time:	15:26:27	Log-Likelihood:	-962.02			
No. Observations:	392	AIC:	1948.			
Df Residuals:	380	BIC:	1996.			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	20.1084	6.696	3.003	0.003	6.943	33.274
cylinders	0.2519	0.326	0.773	0.440	-0.389	0.893
displacement	-0.0169	0.020	-0.828	0.408	-0.057	0.023
horsepower	-0.1635	0.041	-3.971	0.000	-0.244	-0.083
weight	-0.0136	0.003	-5.069	0.000	-0.019	-0.008
acceleration	-2.0884	0.557	-3.752	0.000	-3.183	-0.994
year	0.7810	0.045	17.512	0.000	0.693	0.869
origin	0.6104	0.263	2.320	0.021	0.093	1.128
disp_sq	2.257e-05	3.61e-05	0.626	0.532	-4.83e-05	9.35e-05
horses_sq	0.0004	0.000	2.943	0.003	0.000	0.001
weight_sq	1.514e-06	3.69e-07	4.105	0.000	7.89e-07	2.24e-06
accel_sq	0.0576	0.016	3.496	0.001	0.025	0.090
=====						
Omnibus:	33.614	Durbin-Watson:	1.576			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	77.985			
Skew:	0.438	Prob(JB):	1.16e-17			
Kurtosis:	5.002	Cond. No.	5.13e+08			
=====						

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.

OLS Regression Results

=====			
Dep. Variable:	mpg	R-squared:	0.870
Model:	OLS	Adj. R-squared:	0.866
Method:	Least Squares	F-statistic:	210.7
Date:	Tue, 19 Feb 2019	Prob (F-statistic):	2.25e-159
Time:	15:26:27	Log-Likelihood:	-961.83


```

No. Observations:      392    AIC:      1950.
Df Residuals:         379    BIC:      2001.
Df Model:              12
Covariance Type:      nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          19.8341        6.848        2.896      0.004        6.369       33.299
cylinders        0.2219        0.330        0.673      0.501       -0.427        0.871
displacement    -0.0130        0.021       -0.605      0.546       -0.055        0.029
horsepower     -0.1611        0.041       -3.892      0.000       -0.242       -0.080
weight         -0.0140        0.003       -5.070      0.000       -0.019       -0.009
acceleration    -2.0281        0.566       -3.585      0.000       -3.140       -0.916
year           0.7877        0.046       17.134      0.000        0.697        0.878
Europe          0.9074        0.550        1.650      0.100       -0.174        1.989
Japan           1.2505        0.529        2.365      0.019        0.211        2.290
disp_sq        1.796e-05    3.69e-05        0.487      0.626    -5.45e-05    9.04e-05
horses_sq       0.0004         0.000        2.899      0.004         0.000         0.001
weight_sq       1.554e-06    3.75e-07        4.147      0.000      8.17e-07    2.29e-06
accel_sq        0.0559        0.017        3.349      0.001        0.023        0.089
=====
Omnibus:          32.033    Durbin-Watson:      1.571
Prob(Omnibus):    0.000    Jarque-Bera (JB):    73.327
Skew:             0.420    Prob(JB):            1.19e-16
Kurtosis:         4.945    Cond. No.            5.24e+08
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 5.24e+08. This might indicate that there are strong multicollinearity or other numerical problems.

With origin included as a quantatative variable, the adjusted R^2 statistic is now 0.866, which is better than the 0.818 from our previous regression. The statistical significance of the coefficient on *displacement* has reduced dramatically between the two regressions, and its square's significance is not good either. The statistical signficance of the coefficient for *cylinders* has also fallen dramatically.

With origin included as a categorical variable, the R^2 statistic is now 0.866, which is better than the 0.821 from our previous regression. As with origin as a quantatative variable, the statistical significance of the coefficient on *displacement* has reduced dramatically between the two regressions, and its square's significance is not good either. The statistical signficance of the coefficient for *cylinders* has also fallen dramatically.

1.2.6 Part f

We wish to generated the predicted miles per gallon for a car with 6 cylinders, a displacement of 200, horsepower of 100, a weight of 3100, accleration of 15.1, a model year of 1999, and an origin

of 1 using our regression including the square terms:

```
In [9]: # defining the parameters
const = 1
cylinders = 6
displacement = 200
horsepower = 100
weight = 3100
acceleration = 15.1
# note the 2 digit year, not 4 digit
year = 99
origin = 1
disp_sq = 200**2
horses_sq = 100**2
weight_sq = 3100**2
accel_sq = 15.1**2
europe = 0
japan = 0

predictors_origin = [const, cylinders, displacement, horsepower, weight, acceleration,
                    year, origin, disp_sq, horses_sq, weight_sq, accel_sq]
predictors = [const, cylinders, displacement, horsepower, weight, acceleration, year,
             europe, japan, disp_sq, horses_sq, weight_sq, accel_sq]

print(results2_origin.predict(exog=predictors_origin))
print(results2.predict(exog=predictors))

[38.7321111]
[38.83998021]
```

With origin as a quantitative variable, our model predicts that such a car would get about 38.73 miles per gallon. With origin as a categorical variable, our model predicts that such a car would get about 38.84 miles per gallon.

1.3 Problem 2

1.3.1 Part a

For this, we wish to compute the Euclidean distance between each observation and the origin. For observation 1, this is 3. For observation 2, this is 2. $\sqrt{10}$, about 3.16 $\sqrt{5}$, about 2.23 $\sqrt{2}$, about 1.41 $\sqrt{3}$, about 1.73

1.3.2 Part b

For this, we wish to learn what the KNN prediction for the origin is when $K = 1$. This is Green. As we computed above, the closest observation to the origin is observation 5. Its value is Green. So when $K = 1$, we would classify the origin as Green as well.

1.3.3 Part c

For this, we wish to learn what the KNN prediction for the origin is when $K = 3$. This is Red. As we computed above, the closest three observations are observations 5, 6, and 2. While observation 5 is Green, both Observations 6 and 2 are Red. Thus, when $K = 3$, we would classify the origin as Red.

1.3.4 Part d

If the Bayes optimal decision boundary in the problem is highly non-linear, then we would expect that the best value of K would be ???

1.3.5 Part e

For this, we wish to use Python to estimate the KNN classifier of the test point $X_1 = X_2 = X_3 = 1$ with $K = 2$

```
In [10]: from sklearn import neighbors

# Creating our data
# Observation 7 added as our target, with Green Y randomly
df = pd.DataFrame({"X_1": [0, 2, 0, 0, -1, 1, 1],
                  "X_2": [3, 0, 1, 1, 0, 1, 1],
                  "X_3": [0, 0, 3, 2, 1, 1, 1],
                  "Y": ["R", "R", "R", "G", "G", "R", "G"]},
                  index=[1,2,3,4,5,6,7])

df_train = df[0:6]
X_train = df_train[["X_1", "X_2", "X_3"]]
y_train = df_train["Y"]

df_test = df[6:7]
X_test = df_test[["X_1", "X_2", "X_3"]]
y_test = df_test["Y"]

knn = neighbors.KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train).score(X_test, y_test)
```

```
Out[10]: 0.0
```

For $K = 3$, our dummy test of Green was wrong, so it must be that it is Red. I'm sure that there is a better way to do this, but I was following the code presented in class.

1.4 Problem 3

For this problem, we want to analyze the same auto data as in Problem 1 using a multivariable logistic regression. First, we need to create a binary variable to study.

```
In [11]: # find median of column
median_mpg = autos["mpg"].median()
```

```

# begin new variable at 0
autos["mpg_high"] = 0
# replace values for which mpg is greater than the median
autos.loc[(autos["mpg"] > median_mpg), "mpg_high"] = 1

```

1.4.1 Part a

For this problem, we wish to estimate a logistic regression of our new binary variable on the renamed regressors from Problem 1.

```

In [12]: import numpy as np
import statsmodels.api as sm

# Renaming columns to desired names
autos.rename(columns={"cylinders": "cyl", "displacement": "dspl",
                    "horsepower" : "hpwr", "weight" : "wgt",
                    "acceleration" : "accl", "year" : "yr", "origin" : "orgn"}, inplace=True)

# Dropping na values for logit analysis
autos.dropna(inplace=True)

# Create matrices of X and y values
# Useful when splitting in the next part
X_origin = autos[["cyl", "dspl", "hpwr", "wgt", "accl", "yr", "orgn"]].values
X = autos[["cyl", "dspl", "hpwr", "wgt", "accl", "yr", "Europe", "Japan"]].values
y = autos["mpg_high"].values

# Adding a constant to our X matrix
num_obs = X.shape[0]
const_vec = np.ones(num_obs).reshape((num_obs, 1))
Xconst_origin = np.hstack((const_vec, X_origin))
Xconst = np.hstack((const_vec, X))

# Running the model using statsmodel.api
LogitModel_origin = sm.Logit(y, Xconst_origin)
LogitReg_origin = LogitModel_origin.fit()
print(LogitReg_origin.summary())

LogitModel = sm.Logit(y, Xconst)
LogitReg = LogitModel.fit()
print(LogitReg.summary())

```

Optimization terminated successfully.

Current function value: 0.189320

Iterations 9

Logit Regression Results

```

=====
Dep. Variable:                y    No. Observations:    392

```

```

Model:                Logit    Df Residuals:      384
Method:                MLE      Df Model:          7
Date:                 Tue, 19 Feb 2019    Pseudo R-squ.:    0.7265
Time:                 15:26:29    Log-Likelihood:    -74.213
converged:            True      LL-Null:          -271.30
                                LLR p-value:        4.235e-81

```

	coef	std err	z	P> z	[0.025	0.975]
const	-22.7150	6.140	-3.700	0.000	-34.749	-10.681
x1	-0.0633	0.437	-0.145	0.885	-0.919	0.792
x2	-0.0002	0.013	-0.017	0.987	-0.026	0.025
x3	-0.0399	0.025	-1.618	0.106	-0.088	0.008
x4	-0.0048	0.001	-3.935	0.000	-0.007	-0.002
x5	-0.0178	0.141	-0.126	0.899	-0.294	0.258
x6	0.5196	0.084	6.169	0.000	0.355	0.685
x7	0.4990	0.360	1.385	0.166	-0.207	1.205

Possibly complete quasi-separation: A fraction 0.18 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified. Optimization terminated successfully.

Current function value: 0.183983

Iterations 10

Logit Regression Results

```

Dep. Variable:        y    No. Observations:      392
Model:                Logit    Df Residuals:      383
Method:                MLE      Df Model:          8
Date:                 Tue, 19 Feb 2019    Pseudo R-squ.:    0.7342
Time:                 15:26:29    Log-Likelihood:    -72.121
converged:            True      LL-Null:          -271.30
                                LLR p-value:        4.205e-81

```

	coef	std err	z	P> z	[0.025	0.975]
const	-24.6273	6.285	-3.919	0.000	-36.945	-12.310
x1	-0.1998	0.456	-0.438	0.661	-1.094	0.695
x2	0.0126	0.015	0.850	0.396	-0.016	0.042
x3	-0.0394	0.025	-1.565	0.118	-0.089	0.010
x4	-0.0061	0.001	-4.221	0.000	-0.009	-0.003
x5	-0.0211	0.142	-0.149	0.882	-0.300	0.257
x6	0.5770	0.094	6.161	0.000	0.393	0.761
x7	1.7950	0.761	2.357	0.018	0.303	3.287
x8	1.1271	0.719	1.568	0.117	-0.282	2.536

Possibly complete quasi-separation: A fraction 0.19 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

When including origin as a quantitative variable, we find that β_0 , β_4 , and β_6 , which are the coefficients for the constant, the weight, and the year, are statistically significant at the 5 percent level.

When including origin as a categorical variable, we find that β_0 , β_4 , β_6 , and β_7 , which are the coefficients for the constant, the weight, the year, and being from Japan, are statistically significant at the 5 percent level.

1.4.2 Part b

We wish to randomly and equally divide the data in a training set and a test set.

```
In [13]: # It seems that train_test_split is is model_selection, not cross_validation
from sklearn.model_selection import train_test_split
# following from the problem set
X_train, X_test, y_train, y_test = \
    train_test_split(X, y, test_size = 0.5, random_state = 10)
X_train_origin, X_test_origin, y_train_origin, y_test_origin = \
    train_test_split(X_origin, y, test_size = 0.5, random_state = 10)
```

1.4.3 Part c

We wish to estimate a logistic regression on the training set using the given method

```
In [14]: from sklearn.linear_model import LogisticRegression

# Define the logistic regression
LogReg_origin = LogisticRegression(random_state=0, solver="lbfgs", max_iter=1000)
train_origin = LogReg_origin.fit(X_train_origin, y_train_origin)
print("With origin as a quantitative variable:")
print("Intercept:", train_origin.intercept_)
print("Betas 1 through 6", train_origin.coef_[0,0:6])
print("Beta 7", train_origin.coef_[0,6:7])

LogReg = LogisticRegression(random_state=0, solver="lbfgs", max_iter=1000)
train = LogReg.fit(X_train, y_train)
print("With origin as a categorical variable:")
print("Intercept:", train.intercept_)
print("Betas 1 through 6", train.coef_[0,0:6])
print("Betas 7 and 8", train.coef_[0,6:8])
```

With origin as a quantitative variable:

Intercept: [-30.29184382]

Betas 1 through 6 [-0.99645974 0.02130091 0.01681134 -0.00809184 0.14283266 0.65319314]

Beta 7 [0.41513925]

With origin as a categorical variable:

Intercept: [-29.77528723]

Betas 1 through 6 [-0.95411633 0.02123185 0.01677392 -0.00829327 0.1353172 0.65853586]

Betas 7 and 8 [0.66594604 0.33380007]

1.4.4 Part d

We wish to create predicted values for our test data using our training set and calculate a confusion matrix and classification report.

```
In [15]: from sklearn.metrics import confusion_matrix
         from sklearn.metrics import classification_report

         # Predict new values from our logistic regressions
         y_pred_origin = LogReg_origin.predict(X_test_origin)
         y_pred = LogReg.predict(X_test)

         confusion_matrix_origin = confusion_matrix(y_test_origin, y_pred_origin)
         print(confusion_matrix_origin)
         print(classification_report(y_test_origin, y_pred_origin))

         confusion_matrix = confusion_matrix(y_test, y_pred)
         print(confusion_matrix)
         print(classification_report(y_test, y_pred))
```

```
[[91 14]
```

```
[ 7 84]]
```

	precision	recall	f1-score	support
0	0.93	0.87	0.90	105
1	0.86	0.92	0.89	91
micro avg	0.89	0.89	0.89	196
macro avg	0.89	0.89	0.89	196
weighted avg	0.90	0.89	0.89	196

```
[[90 15]
```

```
[ 7 84]]
```

	precision	recall	f1-score	support
0	0.93	0.86	0.89	105
1	0.85	0.92	0.88	91
micro avg	0.89	0.89	0.89	196
macro avg	0.89	0.89	0.89	196
weighted avg	0.89	0.89	0.89	196

When including origin as a quantitative variable, we are able to correctly classify 91 out of 105 low-mpg cars, and 84 out of 91 high-mpg cars. Respectively these shares are 87 and 92 percent. Of the 98 cars that we classify as low-mpg cars, only 91 actually are. Of the 108 cars we classify as high-mpg cars, only 84 actually are. Respectively, these shares are 93 and 86 percent. Using the average of these, as reflected in the f1-score column of the classification report, we might conclude that our model is better at classifying low-mpg cars than high-mpg cars.

When including origin as a categorical variable, we are able to correctly classify 90 out of 105 low-mpg cars, and 84 out of 91 high-mpg cars. Respectively these shares are 86 and 92 percent. Of the 97 cars that we classify as low-mpg cars, only 90 actually are. Of the 109 cars we classify as high-mpg cars, only 84 actually are. Respectively, these shares are 93 and 85 percent. Using the average of these, as reflected in the f1-score column of the classification report, we might conclude that our model is better at classifying low-mpg cars than high-mpg cars.